

Paper presentation

EfficientNetV2: Smaller Models and Faster Training

Authors :

Mingxing Tan, Quoc V. Le
2021



Content

1. Introduction
2. EfficientNet's Upgrading
3. EfficientNetV2's Architecture Design
4. Main Results
5. Conclusion

Introduction

- The previous Version : EfficientNet (V1) :

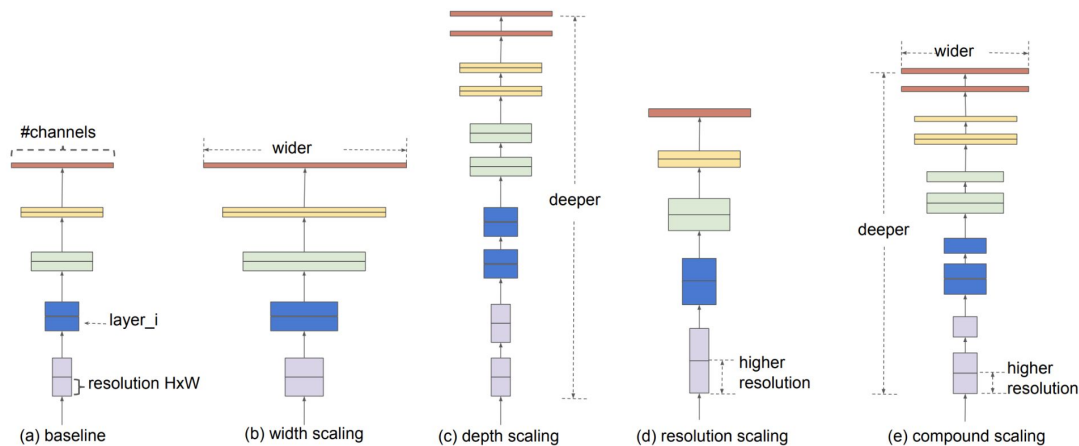


Figure 2. Model Scaling. (a) is a baseline network example; (b)-(d) are conventional scaling that only increases one dimension of network width, depth, or resolution. (e) is our proposed compound scaling method that uniformly scales all three dimensions with a fixed ratio.

Source : EfficientNet, <https://arxiv.org/pdf/1905.11946v5.pdf>

EfficientNet's Upgrading

EfficientNetV1 Drawbacks :


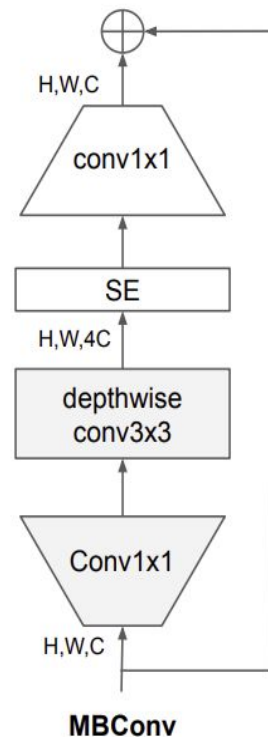
- Training with large images is slow
- Depthwise convolutions are expensive and slow in the early layers
- Compound scaling equally scales up everything

EfficientNet's Upgrading

EfficientNetV1 Drawbacks :

- Depthwise convolutions are slow

	Params (M)	FLOPs (B)	Top-1 Acc.	TPU imgs/sec/core	V100 imgs/sec/gpu
No fused	19.3	4.5	82.8%	262	155
Fused stage1-3	20.0	7.5	83.1%	362	216
Fused stage1-5	43.4	21.3	83.1%	327	223
Fused stage1-7	132.0	34.4	81.7%	254	206



EfficientNet's Upgrading

depth: $d = \alpha^\phi$

width: $w = \beta^\phi$

resolution: $r = \gamma^\phi$

s.t. $\alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$

$\alpha \geq 1, \beta \geq 1, \gamma \geq 1$

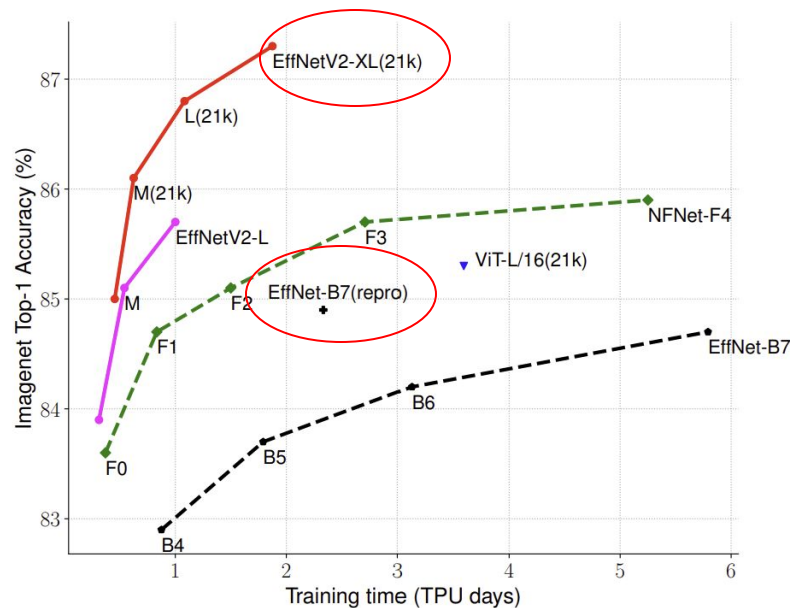
EfficientNet's Upgrading

- New feature in EfficientNetV2 compared to previous version :
 - Progressively increase of the images size during training
 - Progressive learning : adjusting of the regularization
 - Smaller, faster model

EfficientNet's Upgrading

Main goal of this model :

- Reduce training time
- Reduce number of parameter
- While not losing accuracy



EfficientNet's Upgrading

Less parameter and a better accuracy

	EfficientNet (2019)	ResNet-RS (2021)	DeiT/ViT (2021)	EfficientNetV2 (ours)
Top-1 Acc.	84.3%	84.0%	83.1%	83.9%
Parameters	43M	164M	86M	24M

(b) Parameter efficiency.

EfficientNetV2's Architecture Design

EfficientNetV2's Architecture Design

NAS : Training-aware neural architecture search

We need better accuracy (**A**), less training step time (**S**) at the same time, a network with less parameters (**P**).

$$A \cdot S^w \cdot P^v$$

w and v are experimentally determined.

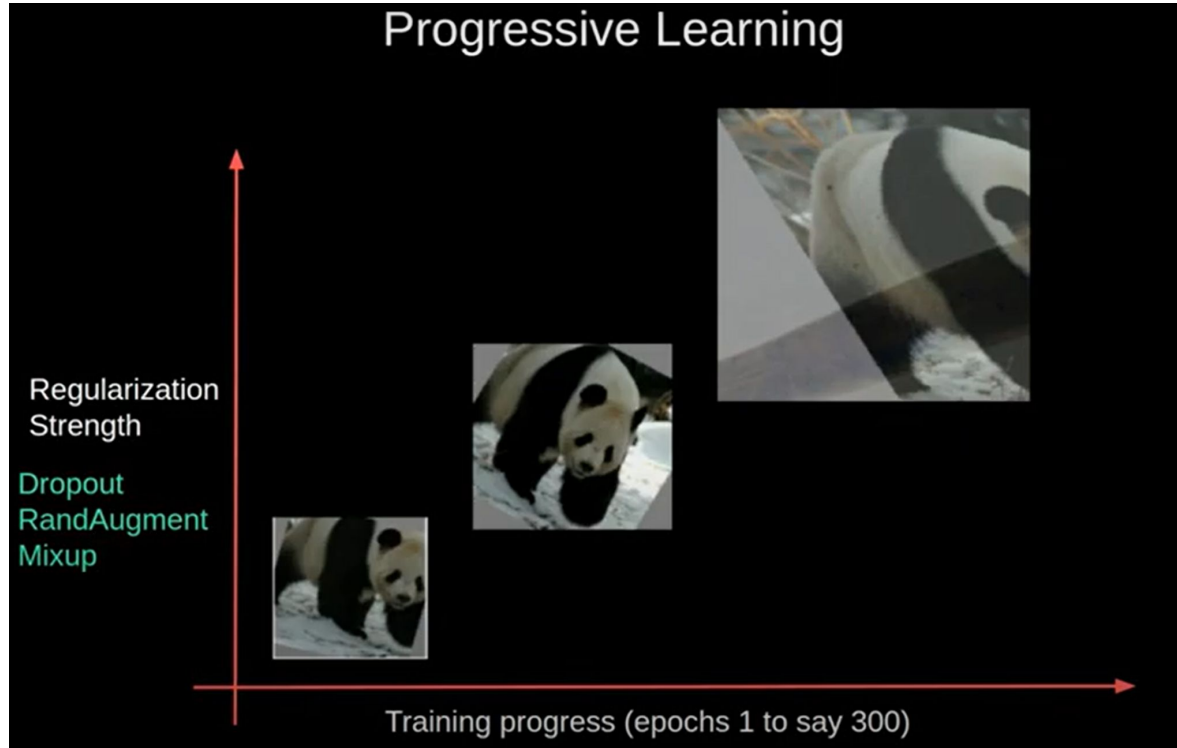
The optimised ones were $w = -0.07$ and $v = -0.05$

EfficientNetV2's Architecture Design

Table 4. EfficientNetV2-S architecture – MBConv and Fused-MBConv blocks are described in Figure 2.

Stage	Operator	Stride	#Channels	#Layers
0	Conv3x3	2	24	1
1	Fused-MBConv1, k3x3	1	24	2
2	Fused-MBConv4, k3x3	2	48	4
3	Fused-MBConv4, k3x3	2	64	4
4	MBConv4, k3x3, SE0.25	2	128	6
5	MBConv6, k3x3, SE0.25	1	160	9
6	MBConv6, k3x3, SE0.25	2	256	15
7	Conv1x1 & Pooling & FC	-	1280	1

EfficientNetV2's Architecture Design



Source : AI Bites, https://www.youtube.com/watch?v=qoSKbMbf1Pw&ab_channel=AlBites

EfficientNetV2's Architecture Design

Progressive Learning : Adapt image size step by step

- **Dropout** (Srivastava et al., 2014): a network-level regularization, which reduces co-adaptation by randomly dropping channels. We will adjust the dropout rate γ .
- **RandAugment** (Cubuk et al., 2020): a per-image data augmentation, with adjustable magnitude ϵ .
- **Mixup** (Zhang et al., 2018): a cross-image data augmentation. Given two images with labels (x_i, y_i) and (x_j, y_j) , it combines them with mixup ratio λ : $\tilde{x}_i = \lambda x_j + (1 - \lambda)x_i$ and $\tilde{y}_i = \lambda y_j + (1 - \lambda)y_i$. We would adjust mixup ratio λ during training.

EfficientNetV2's Architecture Design

Algorithm 1 Progressive learning with adaptive regularization.

Input: Initial image size S_0 and regularization $\{\phi_0^k\}$.

Input: Final image size S_e and regularization $\{\phi_e^k\}$.

Input: Number of total training steps N and stages M .

for $i = 0$ **to** $M - 1$ **do**

Image size: $S_i \leftarrow S_0 + (S_e - S_0) \cdot \frac{i}{M-1}$

Regularization: $R_i \leftarrow \{\phi_i^k = \phi_0^k + (\phi_e^k - \phi_0^k) \cdot \frac{i}{M-1}\}$


Train the model for $\frac{N}{M}$ steps with S_i and R_i .

end for

EfficientNetV2's Architecture Design

Loss function :

- Softmax Cross Entropy :



- L2 Loss :
$$L2LossFunction = \sum_{i=1}^n (y_{true} - y_{predicted})^2$$

Metric :

- Top-1 Accuracy

Main Results

Train dataset :

- ImageNet ILSVRC2012
- ImageNet21
- CIFAR-10
- CIFAR-100
- Flowers
- Cars

Main Results

Table 10. Comparison with the same training settings – Our new EfficientNetV2-M runs faster with less parameters.

		Acc. (%)	Params (M)	FLOPs (B)	TrainTime (h)	InferTime (ms)
V1-B7		85.0	66	38	54	170
V2-M (ours)		85.1	55 (-17%)	24 (-37%)	13 (-76%)	57 (-66%)

Conclusion

Goal Achieved :

- Reduce Model Complexity
- Reduce Training Time
- Outperform previous model

Conclusion

Bibliography and related work

- Tan, M., & Le, Q. V. (2019). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. arXiv preprint arXiv:1905.11946.
<https://arxiv.org/pdf/1905.11946.pdf>
- Tan, M., & Le, Q. V. (2021). EfficientNetV2: Smaller Models and Faster Training. arXiv preprint arXiv:2104.00298.
<https://arxiv.org/pdf/2104.00298.pdf>
- Gordić, A. (2021, May 26). EfficientNetV2 Explained (Smaller Models and Faster Training). YouTube.
https://www.youtube.com/watch?v=CTsSrOKSPNo&ab_channel=AleksaGordi%C4%87-TheAIEpiphany
- Zhang, H., Cisse, M., Dauphin, Y. N., & Lopez-Paz, D. (2018). Mixup: Beyond Empirical Risk Minimization. arXiv preprint arXiv:1710.09412.
<https://arxiv.org/pdf/1710.09412.pdf>
- Zoph, B., & Le, Q. V. (2016). Neural architecture search with reinforcement learning. arXiv preprint arXiv:1611.01578.
<https://arxiv.org/pdf/1611.01578.pdf>


Paper presentation : EfficientNetV2: Smaller Models and Faster Training

Thank you for listening !



Convolutional Layer :


- Perform convolution operation on all the input channel with each filter



Diane Lingrand, CNN:Convolutional Neural Networks Course

Depthwise Convolutional Layer :

- performs a convolution operation on each input channel separately using a different filter for each channel
- Does not mix information between channels
- Prevent overfitting



REMARQUE APRES
PRESENTATION DANS LES
COMMENTS :